# Contents

# Chapter 1

## *The context: Mondeca, the product and my responsibilities*

This report concerns the six months I spent as intern at Mondeca, a French company located in Paris which sells software products capable of exploiting semantic technologies. In particular, their main product is the so-called Smart Content Factory (SCF), a complex system that allows companies to index, annotate and browse through their documentation. Indeed, main targets of the product are those companies which need to handle huge amounts of documents (newspapers, legal offices, and so on).

Let's say that National Geographic (who, by the way, is among Mondeca's customers) wants an automatic way to handle all the scientific articles it publishes onto its website; this is what it will happen:

- it can maintain an ontology with the system
- it may feed SCF with its documents and articles
- SCF analyzes the documents via text mining tools and content annotation systems in order to recognize entities inside the text
- the analysis outputs entities detected in the text, entities that have been inferred from the text (but that are not mentioned) and candidates that could be inserted in the ontology
- a further rule-based classification is possible, meaning SCF will tag the documents according to some rules provided by the customer. For instance, SCF could tag as "Not for kids" the articles that are about crime, blood and sex
- finally, it is possible to browse the whole documentation by means of a semantic search engine

## 1.1 The CAM application

What I will be working on is the so-called Content Annotation Manager (CAM), which is a component of the SCF system. It takes care of text mining a document, annotating it according to the selected ontology and yielding out the enriched document. As I enter the company, CAM is just a step of the process, it is automatically executed and it doesn't have a UI. Data flows in, data flows out. The only web UI that Mondeca has

for CAM is a poorly presented testing environment ("demo") that they demonstrate to potential customers to show off the potential of their API (which ultimately exposes all the goodness the system produces). Thus, my role inside Mondeca would be rethinking the current demo UI in order to provide an easy-to-use, sell-able and pretty web application that could serve the purpose of both convincing potential clients to trust the system and allowing them to actually monitor and intervene over this step of the pipeline.

Here's how it looks like at the moment:



Figure 1.1: Old CAM, Annotations

As it is now, the CAM UI allows the user to do the following:

- select resources to be analyzed
- change some configuration settings
- launch the process
- monitor the results
- add and remove annotations
- read the log messages of the system
- read the RDF file that contains all the results of the analysis

This is a fairly basic overview of what it can be done but, as I undertake the challenge of improving the way the user can interact with the system, I'd rather understand the 10.000 feet view of the product than diving immediately deep into the details and the nuances of it. Such a first analysis gives the right vision of what needs to be focused on,

Figure 1.2: Old CAM, Overview

and, most importantly, allows for a clearer conception of what the underlying servers are capable of producing or doing, in response to eventual user's requests.

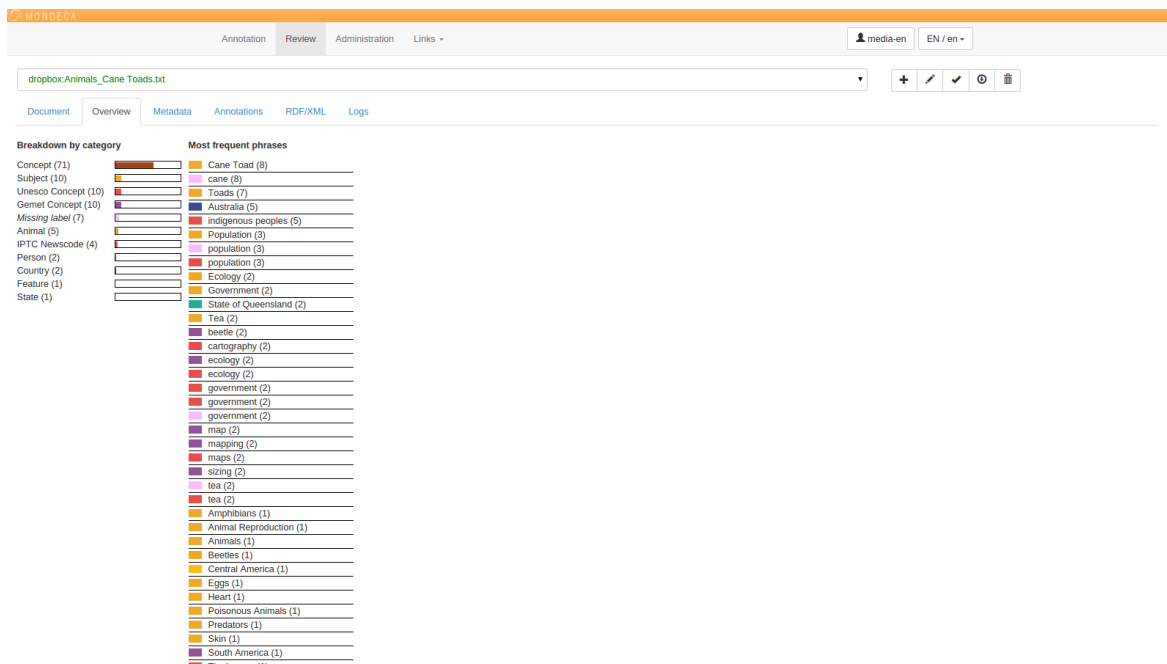However, the main focus of CAM was better defined during the course of the iterative process of sketching and validating the ideas with the team members: while it is mostly used as demo for potential clients at the moment, the idea is to migrate to a fully fledged app, targeting those human agents or indexers who are hourly paid to verify and review automatically assigned tags, bringing on the table their own expertise on a topic. For instance, a surgeon gets paid to delete wrong tags, add missing ones and review the general correctness of the analysis of all the tags related to surgery or body anatomy. The goal becomes then to help such a kind of users to do their job in as little time as possible; the challenge is to demonstrate to potential clients that, through the use of the tools this application provides, human indexers can analyze many more documents in a given time span, than they would by manually reading and annotating the text.

## 1.2 Learning about competitors

During my first days I searched for competing text mining tools with the purpose of focusing on both the different features they present and the ways they present their results, in order to take inspiration in terms of visualization techniques. Such an activity led me to narrow down a list of very simplified features competitors implement:

- Sentiment analysis on entities or keywords
- Sentiment comparisons of a term over different data sources
- Language recognition (e.g. 15% English 85% French)
- Part of Speech tagging
- NER
- NED is rarely advertised
- Relevant words occurrence
- Topics/themes (ordered by pertinence)
- Names translation
- Categorization, often times starting from the bunch of documents only, without prior knowledge
- Adult content detection
- Documents similarity detection

On the visualization side, though, this exercise hasn't been enlightening: on one hand, existing tools often targets really specific markets where presentation details aren't taken too much in consideration, thus the teams working on these solutions focus primarily on functionalities and "smartness" of the software; on the other hand, data visualization is a highly data-dependent discipline, meaning it is ultimately based on what type of data one is willing to show and, most importantly, what kind of reasoning it is trying to support (e.g. - metrics comparison versus content distribution). For these reasons, I decided to go on framing the problem CAM is trying to solve for its users, then to attack eventual visualization challenges from a completely user-centered perspective, meaning my goal would be finding the least complex way to guide the user through the possibly big amount of data the system is capable of producing.

## 1.3   Technical specs

As it becomes more clear what the scope of the project is, it is common practice to delimit the work to be done with a series of technical constraints. In my case, the specifications are pretty loose:

- on the *front-end*, HTML, CSS and JavaScript are the technologies of choice; the running environment is supposed to include every major browser up to the last two versions (Mozilla Firefox, Google Chrome, Internet Explorer and Apple Safari), with a minimum supported resolution of 1024x768. No particular framework or technology is enforced.
- on the *back-end*, the existing code base is written in Java, so that is the language of choice. The system already relies on Shiro as authentication and security framework, and on JBoss as application server.

# Chapter 2

## *The process: technologies, specs and sketches*

In my second month as intern, I went on implementing the initial sketches that served mainly the purpose of better explaining the ideas I had for the new version of the CAM application.

## 2.1 Setting up the project

As it was decided to proceed developing the first two screens, a better definition of the process to be followed had to be stated. Since I was working alone on this project, it was decided to slightly change the way we iterate over the screens, in order to allow me to move both more quickly and more freely. First of all, after having roughly sketched the global flow of the application, I should focus on one screen at a time, completing and validating as many features as possible, before skipping to the next section. In addition to this, since no other developer is involved, we decided to reduce the sketching phase, with the goal of producing semi-working mock ups which are easier to discuss over for non-technical people. I then adopted the following process:

- mock up in the browser a rough representation of a screen
- meet together with the Product Manager and the Sales Manager in order to define features and content
- iterate until the screen is finished
- skip to the next screen

## 2.2 Choosing the framework

The first thing to do at this point, was deciding which framework to adopt so that the development of the functionalities, that enable the application to offer an interactive experience to the user, will be as fast and productive as possible. Naturally, an MVC pattern (or similar) needed to be adopted, as really often happens in web applications and UI development. I decided to go for Angular, a JavaScript framework targeting single page applications, after having considered Backbone and React, as they are

the main players in the front-end frameworks game. While Backbone offers a much more lightweight JavaScript file and great extensibility and freedom in the choice of plugins, template engine, and so on, this comes at the cost of having to manually write a lot of code and browsing through sparse documentation (each plugin has its own). On the other hand React, the new kid in town, developed by the Facebook's team and adopted by Instagram.com, is still in its earliest years, meaning there is not that big of a community yet, plus it somehow reinvents the way web applications are developed nowadays, requiring a little bit of effort by people who come from the already-established way of doing things for the web. So why Angular? There are a lot of reasons why the decision fell on this framework: first of all, it is really quick to get going and easy to learn. My first concern in choosing the right environment, though, was going for solutions that will be easily maintained by the company when I will finish my stage; people at Mondeca in fact are prevalently trained in Java programming, and they have in fact always developed web UIs through the use of the Google Web Toolkit, which allows them to cross-compile Java code into JavaScript and HTML/CSS. Thus, I should make it as straightforward as possible for them to go in and modify something when needed. Plus, a big factor was the availability for Mondeca of an outsourcing partner which already takes care of some of the web UI development for them, and which is a huge expert in Angular-based applications. Therefore, the company can rely on a trusted partner in the future in order to handle my project's evolution. However, Angular has a lot of bonus features that are great for CAM: first of all, it offers out-of-the-box two-ways data binding, which is the automatic synchronization of data between the model and view components. The way that Angular implements data-binding lets you treat the model as the single-source-of-truth in the application. The view is a projection of the model at all times. When the model changes, the view reflects the change, and vice versa. This is of course great in many scenarios, but especially for CAM, which is an application that, at its core, takes as input a series of tags, lets the user browse through it and delete some of the tags, then sends as output the outcome of the user's interactions. In other words, the view needs to simply reflect what's in the model at all times and each user's decision can be simply reflected in the model by deleting (or adding) tags to the list. In addition to this, Angular relies on a declarative way of binding actions and data to HTML tags, which makes it great for fast and mock-up-centered iterations, where the design phase is limited and there's the need for a semi-working mock-up as early as possible in order to let the Product and Sales Managers be involved in the process. In fact, since I work alone on the project, I primarily talk with these people, who are neither designers nor developers, so it's really important that I allow them to focus on the functionality of the application, getting rid of static representations of what the final product will look like. A final consideration on the drawbacks of using Angular needs to be made: it is a moderately heavy framework that could suffer from a performance perspective when thousands of bindings are present at the same time; this is not extremely important for the typical user that's being targeted by CAM: indeed, the user is supposed to be on a laptop (no problems related to low-power mobile devices) and it will certainly be willing to wait

for a small initial loading time.

## 2.3 Automate to enhance productivity

It is common knowledge that, in software development, productivity isn't just related to the time one spends writing down code, rather to the speed tasks get done at. In order to improve such a metric, good developers spend some time making their working environment more ergonomic and smooth; it for this reason that I wish to include in the following some parts of my job that are not directly related to CAM, but that allowed me to reduce the burden of side tasks that affect each developer's everyday job.

**Front-end**

On the front-end side, I made heavy use of tooling and automation, as it is standard practice today. I relied on `gulp` as a task runner, rather than using `make`, for it exists a huge catalog of ready-to-use gulp tasks to enhance every single aspect of the developing activity. What this tool allows to do is basically running standalone JavaScript code (in other words, JS code that doesn't need a browser to be run) in a particular sequence through a command-line interface; moreover, thanks to *watch tasks*, I was able to automatically run my sequence at every saved modifications to my source files. In particular, it was taking care of:

- concatenation, vendor-prefixing and minification of CSS files
- concatenation, uglification and linting of JS files
- minification and in-lining of HTML templates
- compression of images

More on these techniques in Chapter 5. In addition to tasks automation, another area in which tools can improve programmers' life is dependency management; I used `bower` to manage JavaScript and CSS libraries, and `npm` to manage NodeJS modules, like `gulp`'s tasks and external tools. It is especially useful to rely on such tools when cooperating with other people, since they can easily set all the tooling up and fetch all dependencies by simply downloading the source and running the install scripts: `npm install && bower install`.

**Back-end code and deployment**

On the back-end, I already had some Maven tasks to run that were coming from the existing project I inherited. So I simply made compilation and deployment on test server a single job, by creating a `Makefile`. However, such an activity takes several seconds (up to 40s on my laptop), which can be a problem, especially when testing small changes, since long waits may interrupt the mental flow. I thus decided to divide deployment step into two parts, that don't need to go together all the time:

- deployment of back-end code, which consists in copying the WAR file into the server
- deployment of front-end code, which I could speed up by an incredible amount by simply treating the WAR archive as a simple RAR archive, and substituting modified files; this made front-end code's deployment as fast as half a second

**Version control systems**

Finally, I set up a versioning system for both front-end and back-end. On the back-end, I kept using `svn` which is the system of choice at Mondeca, and allows developers to synchronize their project with the internal integration system (Jenkins). On the front-end, I decided to use a local `git` repository, in order to benefit of local branching and smarter changes history.

## 2.4   Framing the problem: specs

In order to get a deeper knowledge of the current state of the application, and to start forming an idea of the desired state of the next version of it, I iterated through a series of sketches; this phase is of the uttermost importance since, by showing concrete drawings of what I think it would be great to show and to hide to the team members, I have the possibility to learn about the product itself. Every team member has something to add to the design phase, since everyone of them is expert in a particular side of CAM (current UI, back-end architecture, system's capabilities, gotchas, and so on). Thus, iteration after iteration, I kept putting together all this knowledge and refining my ideas, until a complete picture of where CAM is now and where it's going next was ready. The next fundamental thing to do then is discussing my ideas with both the Product Manager and the Development Team Leader, in order to collect their feedback and finally redacting a draft of the requirements document, so that the developers, the managers and me could agree on a list of feature to implement. This isn't by any means trying to be a finalized and frozen document, but at the contrary it acts as a reference people can refer to while pushing the product forward.
The result of such a cyclic process, can be trusted here as a way to better define the main purpose of the application and its core sections; let's quickly go through it, keeping in mind that not all the listed features made it to the final realization of the application (especially in the troubleshooting and user sections, which got moved out of scope during the course of the project).

### 2.4.1   Main purpose of the application

The purpose of the CAM web application is allowing the user to load one or more resources and visualize the result of the analysis process. The user will need to see an high-level overview, and a detailed view of the analysis, one document at a time. In

addition to this, privileged users will be capable of browsing the log messages and the actual RDF file that comes out of the processing pipeline.

## 2.4.2 Core sections

In order to properly group up the elements of the application, the following sections will be present:

- File selection
- Overview
- Review
- Troubleshooting
- User profile

**File selection**

This section allows the user to select which files she wants to send to the analysis process. A pool of files coming from different sources is created, shown under the form of a list.

Such sources can be

- Manually inserted text
- Dropbox
- Locally stored file (on the client's machine)
- Remotely stored file (via URL)

The list of files presents, for every file, the following information:

- name of the file
- type of the file
- workflow to be used
- source of the file (e.g. - Dropbox)

The user can "check" files that she wants to

- remove from the pool
- associate with a workflow
- run

The user can:

- remove a file from the pool
- check all

- uncheck all
- run all

While the process is running, the user can keep track of the progress. Both per file progress and global progress.

**Overview**

This section looks and feels as a normal dashboard, and it aims at showing the following measures:

- title of the document
- workflow selected
- number of words analyzed
- number of tags generated
- number of seconds needed to process that document
- metadata of the document, them being

  - customer's decision
  - as many as the customer likes
  - as different in type as needed

- composition of the generated tags, broke down by quality, meaning whether they have been

  - detected in the text
  - inferred
  - suggested as candidates
  - generated as rule-based tags

- composition of the generated tags, broke down by taxonomy
- rule-based tags
- main topic the document is about, exploiting a side-effect of a cluster-based scoring algorithm already in use
- top 3 subtopics with metric, if any (e.g. "the document is also partly about HEALTH 20%, FINANCE 10% and ANIMALS 8%")

The UI has a tiled layout, meaning it is composed of rectangles, each one of those containing a measure or a piece of information. Composition-type measures are presented under the form of pie charts.

**Review**

This section allows the user to browse through the entities that have been found or generated. It is logically divided into two main parts: filters and results. In addition

to these two, another feature is new tags addition: the user can add a new tag (that was missed by the system or that's missing in the taxonomy). This may happen in two ways:

- by the add tag button
- by selecting a word in the text and right clicking it

The process of adding the tag unfolds itself in the following steps:

- type in the new tag
- select its class from a intelligently filled list
- suggest the class if it's missing in the taxonomy

**Review: Filters**   Filters allow the user to narrow down the quantity of information that it is displayed in the results section. Filters are:

- score
- class
- quality, being it
    - detected
    - inferred
    - candidates
    - rule-based

Filters are meant to be collapsible into a side bar which disappears on the left-hand side of the screen. The filters bar itself is structured in a way that it allows the different filters to be collapsed or expanded, in a tree-fashioned style. This allows for further addition of filters.

**Review: Results**   The results are simply 2 different views of the same knowledge:

- selected entities in a list
- selected entities in the text

The list of entities shows, for every entry, the following information:

- name
- taxonomy
- occurrences
- score

The text view displays the document itself and, via highlighting or underlining, it allows the user to locate the selected entities in the text itself.

**Troubleshooting**

This section targets privileged users, who need to read through log messages and browse the whole RDF file which contains the result of the analysis. It is therefore subdivided in two sections:

- RDF
- Log

**Troubleshooting: RDF**   The RDF file is split in Valid and Invalid RDF, by means of tabs. Both valid and invalid RDF is further split into three chunks:

- metadata RDF
- knowledge RDF
- occurrences RDF

The user can decide which one she wants to read through buttons.

The RDF code is displayed in a researchable area: the user can search for a particular word, or tag, and go directly to that part of the code. Tags in the code will allow collapsing and expanding actions, as in regular text editors.

**Troubleshooting: Log**   The log messages are displayed in a table containing the following fields:

- module name
- message type (INFO or ERROR)
- referenced entity
- message code
- processing time
- date
- time of the day
- presence of a stacktrace (YES or NO)

In case a stacktrace is present, clicking on "YES" will lead to a view showing:

- relative message's details
- stacktrace

By means of a button, the user can switch between the "table view" displaying all the messages and the "time view", which breaks down modules by time. This allows the user to understand which parts of the pipeline took the most time, and which parts are the most efficient.

**User profile**

This section is still to be defined.

My suggestions for it would be:

- user's permissions

    - can the user access all the files?
    - can the user access all the workflows?
    - can the user access the troubleshooting section?

- create, edit, delete workflows
- change language
- in the future, past sessions

## 2.5 The first mock-up

After agreeing on the screens, I started to mock up in HTML a static version of the overview page, so that it can be seen how it looks like in the browser. Moreover, it allows the Team Leader and me to start thinking of the theming of the application, phase which usually gets skipped in the sketching process. My first day of mocking up saw me transforming this sketch



Figure 2.1: Sketch, Overview

into this HTML page

Figure 2.2: Mockup, Overview

# Chapter 3

## *The core screens*

In the following chapter, I commit to describe the most important screens of the application, trying to convey both the intended use and the reasoning behind those, which led to the conception of the layout. In other words, my intention is to describe the functionalities of the various parts of the application, while focusing on design decisions. The order of the exposition follows the chronological order in which they were validated and developed through the course of the project.

## 3.1 Overview

The first screen I focused on is the "overview" screen, which is the landing page the user is redirected to after the analysis of a document has completed. The goal is to provide the user a quick glimpse of the most important results of the text mining and content enrichment process, while showing measures that could give some hints about what could be a good course of action from that moment on; so for instance, an administrator could look at this screen and be immediately aware of how good the underlying taxonomy is, whereas a human agent can see if there's something obvious to do, like wrong automatic deletions of tags, or irrelevant terms that shouldn't appear in the most relevant terms list. Selecting the right quantity and type of information to display in this screen was tricky, since there was the need of containing everything in a single screen (no scroll allowed) while providing meaningful insights to both administrators (who primarily care about taxonomy's quality and annotating system's performance) and human indexers (who would like to be helped in the process of understanding what happened during the analysis and what can be easily fixed). For these reasons, my design includes just 9 "widgets", organized in 3 rows, so that every piece of information is effortlessly reachable by the eye, being it widely separated by blank space from the other elements of the page. Measures and, more generally, data about the document, are grouped in a logical manner: on the left-hand side, "wordy" descriptive information is presented, such as metadata on the document itself, number of words analyzed, configuration settings applied to the pipeline in order to process this resource (mainly a list of the taxonomies involved and whether or not the content classifier was enabled). On the right-hand side, the real dashboard-like content, laid out top to bottom from most informative to most general. As we can see in the screenshot, in the first row I decided to put

information that is widely related to the aboutness of the document itself, being it a high-level dissection of the main topics it touches, plus a sequence of tags produced by the underlying classifier, which is capable of spotting existing relationships between elements that come from the global knowledge, which it is formed over the document's content: what this means is that, basing on some rules configured by the customer, the system can classify the document putting together both what it has been found in the text, and what it has been inferred by means of the taxonomy.



Figure 3.1: Overview screen, version A

So, in the example in the picture, the classifier was able to determine that a given article can be identified as "Not for kids" basing on the presence of "adult" terminology, content or topic. What it's interesting here, is that the Web API doesn't - yet - provide the list of "main topics", as I called them in my mock-up; the idea of showing this incredibly useful and insightful piece of information came from a discussion that I had with the Sales Manager and one of the core back-end developers of the team. This developer was briefly explaining an algorithm he personally invented to score entities found in the text, by means of clusters. What we realized is that, since the fundamental nature of the model they're using is tree-like, meaning there's always available a broader-narrower relationship between resources that allows to bring the model down to a tree, this algorithm can always represent entities detected in the text and their relationships in a way that allows to split such a tree in clusters. By clustering together different entities under a cluster's root node, the system will easily output relevancy scores for the given terms; thus, in a document presenting a big cluster having as root node the "health" term, and a small cluster having as root "finance", a term belonging the first cluster will be considered as "more relevant" than a term belonging

to the second cluster, by weighing more occurrences of the former than occurrences of the latter. The really interesting notion here is this clustering process, that seems to perform very well and which gives out some interesting analysis of the content of the resource being analyzed; however, clusters are just a step of the whole algorithm, and they are not stored anywhere, nor kept in memory after scoring completes. What my design is suggesting is to revisit the already existing software so that this clusters can be sent to the client application. I strongly believe that this is really useful for CAM, especially when we are thinking of a dashboard, aiming to give the most concise summary of what the annotation process has done, with respect to what's actually present in the text. What follows in the screen is somehow more "standard" statistics: in the second row, I put the total number of words analyzed, which is important in order to correctly interpret what follows, a breakdown of all the tags by type, or category, answering to questions like "how many of the tags that we have are coming from the text?", "how many of those have been inferred by the system?" and so on. Finally, the top 5 most relevant terms, organized in a bar chart. The third and last row contains measures that are not less important, yet they are placed in the bottom area of the screen, since they are somehow less related to the document itself, but rather more to the quality of the system and of the model. In fact, here the user can find a timing estimate of how long the process was, a breakdown by taxonomies that were involved in the pre configured workflow (which complements the lists already shown in the sidebar, since it's very likely that not all those taxonomies actually yielded out some term) and the top 5 most frequent classes. In such a way, administrators or taxonomists can visualize which taxonomy is useful and which is not, which classes are always triggered and which ones are probably not that important in the business model of their company. In conclusion, there is another important thing to notice about this first Overview screen: the whole aboutness row relies on the presence of two plugins, the cluster-based scoring plugin and the rule-based classification plugin. But, as things are now, it is entirely up to the customer to configure those plugins and drop them into their pipeline. What this implies for the design of the screen is that we could not have the first row, since the data it depends on could not be available. The solution for this was designing a "plan B", introducing a new row, that shows three highly correlated charts: the most relevant terms, as it is already shown in the original version, the most frequent terms, which is basically the same chart ordered by occurrences, and the most frequent AND relevant terms, which is a combination of both the previous charts. These three graphics, when compared with one another, tell the user much a deeper story about the terms that are found in the text, thus forming a great fallback in absence of the "aboutness" data.

## 3.2 Review

Regarding the "review" screen, its main goal is to enable the users to quickly browse through a possibly very long list of tags, understand the logic that made the system

Figure 3.2: Overview screen, version B

annotate the document with those tags and, in case she considers some annotation wrong, remove one or more particular tags. Moreover, whenever the user realizes that an important tag is missing, either because it wasn't detected in the text or because she thinks it's correlated and useful to the company owning the content, she can add it: let's say, an article covering some aspect of Michael Schumacher's life, but never mentioning F1, can still be annotated with a "F1" tag by a human agent, in case this couldn't be automatically inferred by the system. With these objectives in mind I approached the design of this screen from the user's point of view entirely. In order to better express how I did this, let's rely on a user story:

> "a human agent expert in some given area is provided with a list of already annotated documents to verify and review, being paid basing on the throughput of documents"

Having a user story, no matter how brief and straightforward it may be, is powerful way to stay focused on what the user's needs and goals are; in fact, since I knew from the beginning what the final experience of my targets should be, I could concentrate on what really matters to them. My idea is to get rid of every static representation of the outcome of the analysis, and to allow people to narrow down the amount of information they are presented with letting them focusing on the possibly few annotations they are interested in. In order to achieve this kind of user experience, I provide two simple yet

powerful tools: a table of tags and some filters. By using the filters, a human indexer can easily see only the tags she is supposed to review and get her job done in the smallest amount of time possible; moreover, by making good use of the high-level information she has just seen in the Overview screen, quick action can be immediately taken, without having to spend too much time digging into the text, details, comparisons, and so on. At this point, I was more than ever decided to go for a single-page web application: it doesn't feel right to have to refresh the page every time the user hits on a filter, as it usually happens in the competitor's' tools; for this reason I feel like adopting a heavy client-side framework like Angular was the right call. The third most important element of a UI of this kind is, of course, the original text, and the context in which detected terms are found: in other word, it is important to have access to the position in the text where an entity lives, so to not only be able to recognize whether or not the annotation was correct, but also to understand why an element appears on the list (it happened in the past of not understanding why there was a "bear" in the list, because the system simply normalized the past participle "born" into "bear"; in such cases, it's extremely important to be able to track back to the analyzed text). In order to address this, I decided to split the screen into two equally sized columns, one containing the table of tags, the other containing the annotated text, while hiding behind a toggle button a drawer panel containing the filters. A further adjustment to the design that I made was not overlaying the panel to the two columns, rather "pushing" the content to the right, basically displaying simultaneously on the page three columns: filters, tags and text. This was done with the goal in mind of allowing the user to see what's changing in the list and in the annotated document, as she filters out entries; and this is made especially important by the interactive nature of the application I set out to build, since it would be useless to have a single-page web application, capable of instantaneously react to user's input, if the filters' bar covered partially the resulting outcome of those actions.

Finally, always along the lines of helping a human agent do her job quickly and promptly, I added breadcrumb faceted navigation on top of the screen, to help the user keep track of what filters are enabled, at any given time. This may seem a strange decision, since there's no actual hierarchy in the filtering process, but it is still useful to have it as a feature in the app, since we don't want CAM's users to validate tags just because they forgot that something can be out of the list that's on the screen in that moment. This also gave me a chance to put a "reset filters" button in a place where it could be intuitive for the user to find it.

## 3.3 Analyze

Introducing the "analyze" screen (in lack of a better name). This screen is probably going to be the first one that the users are going to interact with: indeed, it contains the core controls which enable them to choose the resources to submit to the analysis process, see the progress of said analysis and then dive into the results for the selected
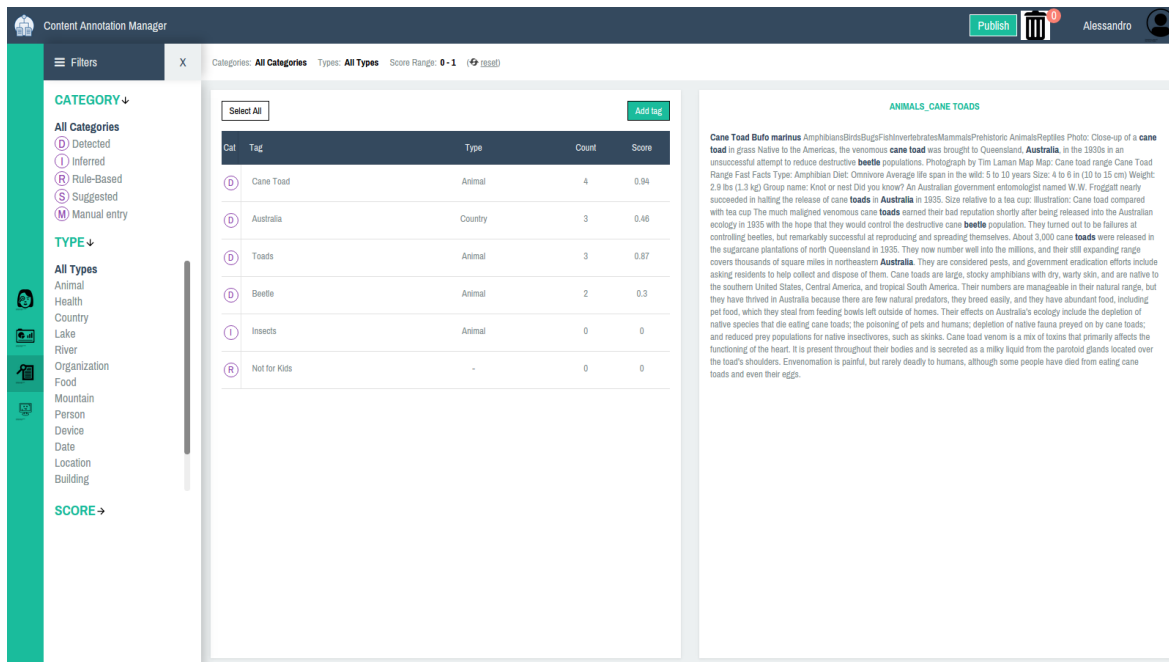
Figure 3.3: Review screen, filters open



Figure 3.4: Review screen, breadcrumb faceted navigation

resource. Moreover, here one can change some of the workflow's parameters: remove some taxonomies from the list of the ones that will be used, disable the classifier or specify which rules should it take into account during its classification process (which is, indeed, rule-based). In order to design a solution for all of these tasks, I decided to split the needs the user might have in two sections: on one hand there is the resource selection, which should include buttons to add files/resources, a list of the selected ones, a way to change the parameters, and so on; on the other hand, there is the process monitoring, which should provide visual feedback for the user, such as an indication of how long the process will take, which resources have already been processed, and so on. After having done this, I tackled the design challenges and, through a sequence of iterations, I came up with a solution composed by the following elements:

A sidebar, on the left-hand side of the page, visually representing the process monitoring section; it includes, from top to bottom: a button to perform the most important action - running the analysis on all the selected resources an overview of the work already done and yet to be the global progress (e.g. - 40% completed) some

statistics on the results of the review of the documents performed by the user itself (e.g. removed tags) The choice of having a sidebar comes from two considerations: first, it is consistent with the general design of the application, which presents a sidebar in almost every section, and secondly, it makes perfect sense to logically separate the "action", displaced in the main area of the page, from the "details". From a design point of view, grouping is the technique I relied upon the most in order to visually represent the relationship existing among the various pieces of data. In fact, by doing this, it becomes obvious that what's going to appear and change on the left, is just an addition, or an extra if you will, to what occupies the main portion of the screen, which instead is the most important thing to pay attention to.



Figure 3.5: Analyze screen, global glance

The main content, visually representing the resource selection section. Here, the most prominent element is the list of selected documents to be indexed; indeed, I decided to represent each resource with a single, well-defined, block, rather than with a simple entry in a table or list.

This is meant to be a visual signal that every single one of those reveals a big process, involving the content augmentation job performed by the underlying system, the tags review and modification by the human agent, while requiring special attention in the configuration and parameters selection, since every workflow (as it is called in CAM) relates to a specific type of content, to some types of file (if the document is contained in a file, as it often happens with CAM), and it has been configured by carefully choosing the plugins to be activated during the analysis (which, for instance, affect the way relevancy scores are computed). Therefore, every document is spaced-out

from the others, has its own status indicator (not processed, ready, reviewed) and its own configuration.

In this section of the application, there is a lot going on. Thus, my approach consists in hiding as much complexity as possible by providing a single point of contact between client and server. Indeed, the main object the user interacts with is a "pool" of documents to be sent to the server; she can populate such a pool in different ways: adding files from a pre-configured Dropbox folder or the local hard drive, by pointing to web resources through URLs, or even by manually entering text (this functionality hasn't been developed yet, but it will be). There are many things the user can do, a part from populating the pool, such as setting different workflows to different documents, running or re-running only some of the resources, filtering the list, changing the parameters characterizing the workflows themselves (e.g. which taxonomies to use), and so on, but, by putting the list of "to-dos" that will be sent to server at the center, not only the user doesn't get overwhelmed by the variety of options and nuances, but - most importantly - she is always under control.

# Chapter 4

## Focusing on the user: from useful to usable

Even though the technology involved is powerful and capable of really interesting things, when designing a User Interface, the main goal should always be focusing on the user, finding in the realization of such design that sweet spot between extremely simple and usable and very effective and useful. Achieving the result is not always immediate and easy, and consists in a process of constantly putting the user at the center. In the following, it is explained how I reasoned about this in the most important page of the application, plus some considerations on how animations can improve the overall UX.

### 4.1    A UCD-based approach: the Review screen

There are a couple of aspects of the review screen that certainly demonstrate how designing for the user, or User Centered Design (UCD), is a great approach to make more usable products. In fact, some small additions and modification to the global behavior of CAM are entirely due to this way of working, like the "reset filters" button already mentioned, which enhances the user experience by reducing the number of clicks to switch from one filter to the other by a great deal, adopting a single color to highlight the terms in the text instead of using a color per class or category, which greatly improves the readability of the text, or providing affordances for the user to undo an action or rollback to a previous state, for example when she deletes some tags.
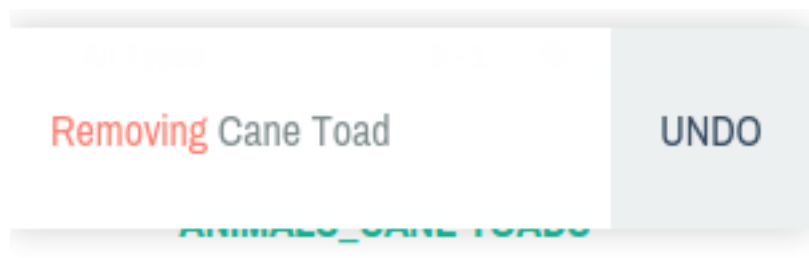


Figure 4.1: Review screen, removing notification, with possibility to undo

The greatest challenge I am facing designing the new version of CAM is being able of both improving the existing demo and suggest new features to be integrated in the application, so that it could be used as a stand-alone tool in the future. While

doing this, I had the chance to modify some aspects of the user's flow, revisiting the experience. A good example of this is the "add tag" action, which used to consist in the steps:

- select Class
- select Label among the ones corresponding to the selected class

This probably makes a lot of sense to a developer who also knows how the taxonomy works, but it doesn't tell much to a human indexer who usually isn't trained in those terms. What actually happens during the usage of this function is:

- the user thinks of a tag she wants to add
- the user tries to help taxonomists by indicating a class

I made the process more user friendly, by guiding through the tag addition with a wizard, giving also the possibility not to suggest any class, since one could not know what to do with the tag she is thinking of. Also, I got rid of the notion of "class" by using the "type" term, which can be less confusing for some users. In addition to this, I improved the way new tags get integrated in the UI, by adding a "manual entry" type, that can be reached through the filters. As a result, the user has a quick method to revisit what she added, and possibly change her mind.

At this point, the "review" screen was still lacking a couple of fundamental features: a way to submit the modifications and store them on the server, and the possibility to find a tag in the document and easily see its details in the table. A usability problem that I spotted in the old version of CAM was the complete lack of a system to undo just executed actions or rollback past decisions (in other words, to change one's mind). For instance, when the user deletes a tag, it simply disappears from the application and there's no way to restore it; or if she adds a couple, they get lost in the midst of the possibly very long list of tags. Finally, by hitting a "done" button, the user would cause an automatic process of storing the modifications, no way back. By analyzing such a situation, I immediately thought of an useful parallel with the shopping carts in e-shopping websites: when you go on Amazon.com, there is no way you can accidently end up buying something you didn't want, nor it happens that you leave an item out of your order. This is mostly due to the last step of the purchasing process, represented by the shopping cart, which summarizes all the details, the items you are about to shop, prices, and so on; in addition, there's some confirmation action required in order to proceed further with the order. The way I see CAM is greatly linked to this scenario: CAM's users decide among a catalogue of tags which ones they want, which ones the would discard, and they can even "buy" some tags that are not present in the catalogue itself. Therefore, my suggestion is to add a simple "summary" step, in which the users are presented with a list of the modifications that are about to be permanently stored, with the option to roll back and change their minds.

Another idea that came up aiming at supporting usability of the application is allowing a "right-to-left" interaction, meaning: instead of browsing the resources from

Figure 4.2: Review screen, summary

left-to-right, starting from the filters, then reading the table, and finally finding the interesting tags in the document, the user might want to be able to quickly spot a tag in the document, and then easily find the tag in the table, without having to scroll a possibly very long list of those. The interaction to be realized is simple: clicking on the highlighted tag in the document should somehow cause a further filtering in the table on the left.

## 4.2   User Experience: Animation as a tool

Contrary to what one may be tempted to believe, animation is not only a way to make things prettier in an interface, but rather it can be a powerful tool in the hands of a good designer: indeed, by making good use of motion techniques and by paying attention on not falling in the trap of cluttering the screen with too many moving elements, it is possible to greatly enhance the usability of a product. As humans, when we interact with some piece of software, we tend to assume that it will behave as real things do in the real word. The term behaviour is pretty powerful, but it helps conveying the message that we implicitly think of shapes and UI components as real things, and as such, we expect them to follow the same basic principles:

- objects accelerate and decelerate, rather than abruptly stopping when they have reached their destination
- objects don't enter in and exit from sight "teleporting" themselves into view

When this implicit pact is broken by a UI element, two things happen in one's mind:

- first, the user needs to realize that something happened; for instance, a dropdown menu changing from collapsed state to expanded state
- then, the user's brain needs to figure out what happened; in this example, it needs to reconstruct all the missing 'frames', in order to figure out the change of state

This process is really fast, but it has the negative effect of interrupting the user's flow and, therefore, it affects the overall experience. Thus, keeping in mind these principles, I played with animations and implemented smooth state changes, with the goal of easing the user's way through the application; indeed, since animated transitions can work as intermediaries between different UI states and can help orientate users, I focused on two major state changes in the "review" screen and studied how I could put all these notions together and help the user understanding not only the sequence of events, but also sorting out causes and effects. In the "review" screen, there is the main content, as usual, and some additional filters to act on in a sidebar. This bar isn't always needed, so it makes sense to hide it behind a button and to show it only when the user really wants to interact with it; however, when visible, it doesn't overlap the rest of the content, but rather it pushes it on the right, causing a big change in all the elements the user just got used to seeing. In presence of big state changes like this, it is important that the UI doesn't surprise the user in a negative way, by moving things on the screen without her understanding where they went, or where they came from. I thus decided to animate the scene, so that the user can follow the motion of an element to understand how the before state of the page and the after state relate to each other.
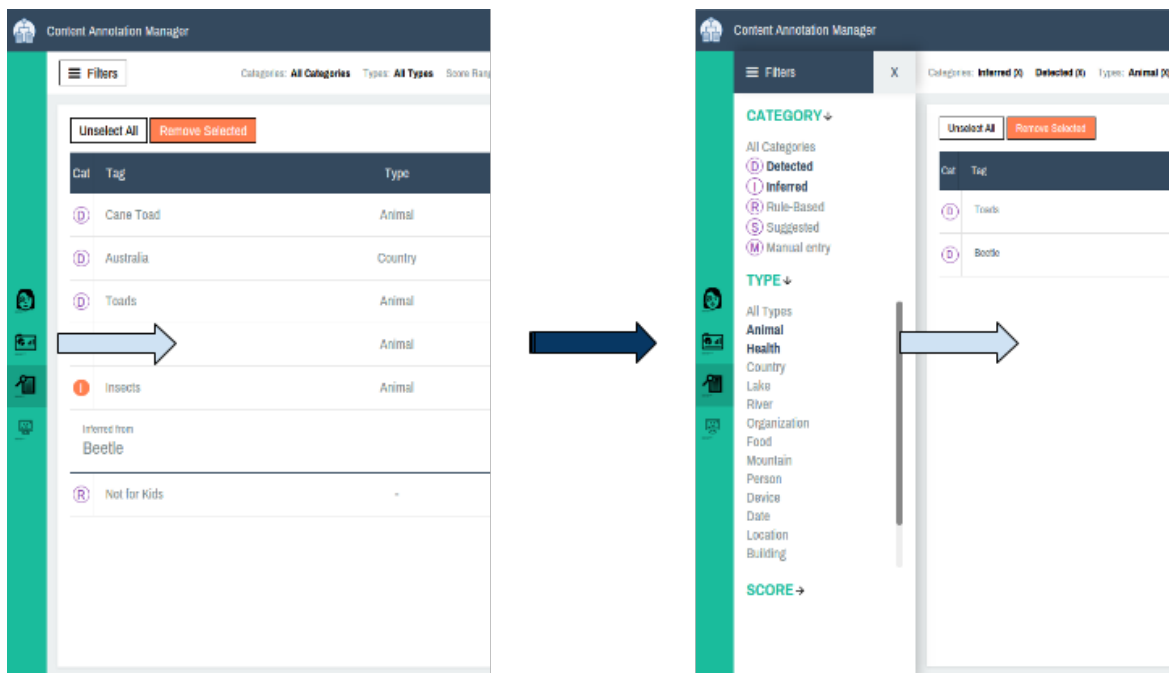


Figure 4.3: Review screen, sidebar animation

Another abrupt change that I wanted to avoid is the elimination of a tag. When a tag is deleted from the table, it disappears and can be found, and possibly restored, by clicking on the bin icon on the upper-right corner of the application. This is a perfect example of how a simple animation can solve a big usability issue, since it takes a huge mental leap in order to understand that an entry in the table just got transferred to another place of the application, without having to remember where things go when they get deleted. Also, this problem shows how true it is that we always expect UI's objects to behave as normal things in the real word: when you throw away something, it doesn't materialize itself in the trash bin, then it shouldn't be happening in the virtual world. So how can an animation solve this issue? Nothing too fancy, but a simple notification-style popup that slides in from the right in the upper-right corner of the screen, accomplishing two tasks at once: on one hand, it gives the user the possibility to undo the action, while on the other hand, it suggests, or reminds, where the tag has just gone to, by sliding up towards the bin icon when disappearing.
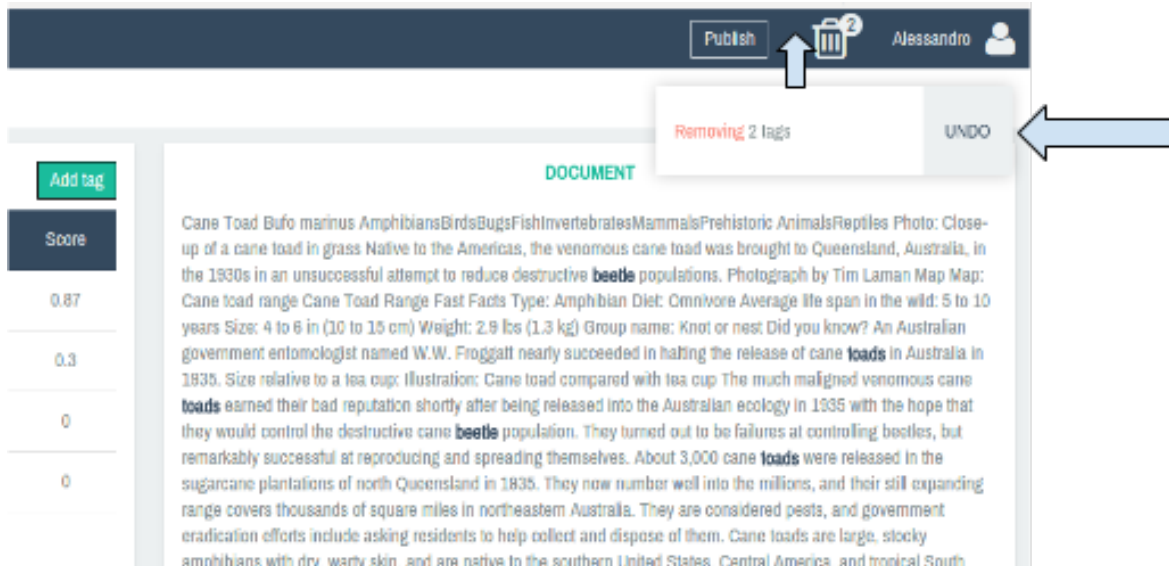


Figure 4.4: Review screen, notification animation

# Chapter 5

## *Perfmatters: a deep dive into performance*

When it comes to performance on the web, more and more often developers don't take into account the external factors that constrain an application's speed, which don't depend on architectural patterns or optimized data structures. However, those are the real bottlenecks in the web of today, especially when targeting a desktop environment, where developers can count on medium to high processing power. I believe this is a really important problem to address for CAM, and other Mondeca's products as well, since we are not producing applications to be used on mobile; in addition, most of the jobs and features that such products offer are run on servers, therefore the client needs to continuously rely on network requests and gets slowed down by the decently big amount of data that needs to be transferred on the wire. For these reasons, it feels compelling to reduce the impact on performance of elements that are independent from the execution of the jobs themselves. Most commonly, such elements are the following:

- *file size*, which basically determines the amount of time spent by the browser on downloading the assets needed to display the requested resource; this can be greatly reduced by means of established techniques such as minification and compression of assets
- *number of HTTP requests*, which greatly affects the bootstrap of a web page since every HTTP connection that the client needs to open adds a fairly significant overhead, which is generally big compared to the time the client will actually spend downloading the resource, especially when such a resource has been minified and compressed; it exists a simple and powerful solution to this problems, which consists in concatenating as many resources as possible in a single file. This technique, when combined with the previous ones that aim at reducing the file size, can optimize the loading time by a great factor

By means of concatenation, minification and compression techniques, the page can load faster and feel instantaneous to the user. However, other things can be done in order to further improve the experience for the user; first of all, though, it's important to have valid measurements in order to understand what needs to be improved, take action and then validate the results. Thus, I performed a performance audit on CAM, tested some solutions and then evaluated the effects against some known indexes and thresholds.

## 5.1   Two metrics: loading speed

As the current state of the application is at its earliest stages, I mainly focused on two "metrics" for performance gage: loading speed and runtime smoothness. Please also note that, at the time these measurements were taken, no actual server-side processing was taking place, and the actual data shown by the application was mocked-up and dished out as a JSON file to the requesting client. Even though this may seem to be invalidating the measurements, being that I am referring to a situation that it's not going to be replicated in real use case scenarios, it is actually meaningful to try and optimize the performance of the application at this point in time, since my goal, as a front-end developer, is to reduce the impact of the application shell on the global performance as much as possible, possibly shooting at adding an overhead to the whole process of indexing documents that can be considered as negligible.

Let's start with loading speed. The most common metric referred to when assessing loading speed is the so called speed index, which basically is the average time at which visible parts of the page are displayed, expressed in milliseconds. As the *webpagetest.org* documentation puts it (emphasis added)

> "The Speed Index metric was added to WebPagetest in April, 2012 and measures **how quickly the page contents are visually populated** (where lower numbers are better). It is particularly useful for comparing experiences of pages against each other (before/after optimizing, my site vs competitor, etc) and should be used in combination with the other metrics (load time, start render, etc) to better understand a site's performance."

Webpagetest.org is one of the most popular and useful online tools for performance evaluation on the web; it basically runs a tests suite over a given URL and outputs a series of tips to improve the performance and important measures to be aware of. For the sake of this first experiment, I am only reporting the Speed Index, which was **4575**. This figure doesn't tell a lot per se, but it is important to know that a solid speed index value is around 1000, as it will be covered in the next sections.

In order to dig deeper into what could be causing a lower loading speed than expected, I took a look at both the network usage and the JavaScript profile, thanks to the Chrome DevTools. The network traffic analysis revealed that the client was spending too much time downloading the resources, and this was slowing down the first paint of the page; thus, the first improvement I made was relying on the established techniques explained in the previous paragraphs: indeed, by means of concatenation, minification and compression techniques the time to first paint dropped down from roughly 1 second to 0.5 seconds, on my testing environment. The reason why this happens lies in the way the browser's parser is dealing with external resources when first analyzing (and then rendering) a page: it goes from top to bottom through the

HTML file, it fetches the resource from the server every time it encounters a script tag or style link, and only then, the first paint appearing on screen can take place. It needs to be acknowledged that the speed index is taking into account the time at which the page is **mostly visually complete**, whereas the time to first paint it's just the moment at which the browser can actually **start rendering** pixels on the screen. There is an important difference between the two metrics, which can be brought down to a single fundamental concept: perceived performance.

## 5.2   An orthogonal metric: perceived performance

At the same way UCD predicates designing for the user, we shall put the user at the center while improving performance. Opposed to computational performance, perceived performance refers to how quickly a process appears to be completed to the user, which may often differ from the actual speed that process has been completed at. For instance, the amount of time an application takes to start up, or a file to download, is not made any faster by showing a splash screen or a progress bar. However, it satisfies some human needs: it appears faster to the user as well as providing a visual cue to let them know the system is handling their request. This is a very relied upon technique when trying to deliver a great experience of use for a not-so-snappy system; when it comes to web applications, the goal is to be able to display parts or previews of the final status the screen will reach, while loading it. Therefore, coming back to difference between speed index and time to first paint, the latter doesn't give an accurate indication of how responsive the application will be **perceived**. In this context, recently it has been gaining a lot of popularity a new technique consisting in showing the application shell as quickly as possible, then dynamically injecting content into the page. In order to achieve this, I had to first solve another issue that was heavily affecting the pages load time, causing that 4000+ speed index to be so high; that is, an Angular project is typically composed of a lot of HTML templates, in order to enhance code readability and maintainability. However, the downside of such a module-based approach, is that every directive or template included in the page is going to trigger an HTTP GET to the server requesting for the corresponding HTML file. So how do we solve this problem? I tested three solutions, basically building every new idea on top of the previous one.

At first, I enabled templates **caching** through use of Angular's `$templateCache` service. This allows Angular to fetch every template only once, and reuse it throughout the whole application every time it is requested again. While such a strategy improves the overall navigation experience, it still doesn't solve the problem of having to contact the server the first time an asset is needed. Then, I instructed the application to **prefetch** all the HTML assets after all the core requests have been handled and the first page has been completely loaded. The objective of such a strategy is to prepare the cache with all the assets that will be needed in other sections or pages, even they may be hidden at first. While this two improvements, when applied together, greatly reduce the delays introduced by the need of continuously fetching the HTML templates,

they don't improve the perceived speed of the application, rather they deal with the computational one. The user was still looking at a blank page until all the resources were ready: indeed, such caching and prefetching strategies rely on Angular to be fully loaded and functioning. Therefore, I decided to **inline** all the templates into a single, minified, JavaScript file (under the form of simple strings) and to include such a resource into the index.html. But why is this any better than the previous solutions? Here's what happens in the browser in the two situations:

Without inlining:

- the browser fetches the index.html
- the browser fetches the script.min.js, which is the file containing all the application's code
- the browser runs such script, which will then fetch the missing templates from the server
- while the network request is still pending, the browser goes on parsing and rendering the first page, asking **to the server** every HTML template it needs for the first load

With inlining:

- the browser fetches the index.html
- the browser fetches the script.min.js
- the browser fetches a brand new templates.min.js file, containing inline declarations of all the templates used throughout the application
- the browser goes on parsing and rendering the first page, asking **to the cache** every HTML template it needs for the first load

Even though the difference may not seem that big, it has great implications on the perceived loading speed. Indeed, the user is now almost immediately shown the application shell, meaning that the whole structure of the application (header, navigation bar, sidebar, and so on) is almost immediately rendered, while the dynamically injected content gets displayed as soon as it is available. However, in order to better define the viability of this approach, one must define what "almost immediately" means, and understand why adding one HTTP request at the beginning of the page loading process (that template.min.js additional resource) feels faster to the user, than it does fetching an HTML template when needed. The "application shell" strategy is backed up by several studies, and has its roots in the publication *Response Times: The 3 Important Limits* by Jakob Nielsen on January 1, 1993. Thanks to this study, we can now quantify qualitative measurements such as "immediate" and "instantaneous": indeed, Nielsen proved that a waiting time under 100ms feels instantaneous, while a waiting time around 1000ms signifies a context switch to the user. Further analysis have further expanded these concepts and, most recently, have led to the definition of the RAIL model, which is defined as by the Chrome team who came up with this new concept as (emphasis mine):

> "a **user-centric performance model**. Every web app has these four distinct aspects to its life cycle, and performance fits into them in very different ways: Response, Animation, Idle and Load"

Basically, the RAIL model gives us developers a few figures that we can refer to whenever we are facing challenges in improving the perceived performance; indeed, thanks to RAIL, we are finally left with the following timing constraints:

- 16ms per frame during an animation
- 100ms of response time, upon user's interaction
- 1000ms of loading time

But let's bring this back to CAM. This model explains why the inlining strategy is so successful: during load time, we're working with a 1000ms time span, during which the browser can fetch the additional templates.min.js file. It's perfectly fine for the user to be waiting up to 1 second for the application to load, thus one can make use of this and load the templates resource; as a consequence, two milestones are achieved:

- the user sees the application shell all at once, in under a second time, and mentally establishes that the loading phase is done (even if it's not);
- when the user will change views or pages, all the structural element will be immediately in place, giving the perception of instantaneous navigation

Here's what the pages look like before the content is dynamically injected.

## 5.3 Two metrics: runtime smoothness

Regarding the runtime smoothness of the application, a few tweaks needed to be done. First of all, the meaningful metric here is the frames per seconds (FPS), the goal is to ensure that animations and scrolling play out at 60fps, and again I relied on the Chrome DevTools in order to reliably get this information. Since the "analyze" screen is basically mostly static, I focused on the "review" and "overview" screen; both of them used to have some performance issues in terms of animation smoothness, but the origin of such low fps turned out to be very different in the two cases. In the "overview" screen, a visible lag could be seen while first loading the page, and this feeling was confirmed by the depressing value of 3fps in some steps; indeed, the charts are animated (bars grow to their final size, pies build up in a fancy way) as soon as the page is loaded. A deep look into the Timeline view in the tools revealed what was causing it: the Angular's directive I am using to stamp out those charts (which is basically wrapping the Highcharts.js API) was running its underlying JavaScript in order to make the animation happen. Regardless of whether or not this might be the right way to do

Figure 5.1: Overview screen, application shell



Figure 5.2: Review screen, application shell

it (more often than not, such JavaScript-empowered transitions do not rely on GPU acceleration, even though I can't be sure since I didn't dig into the implementation code of it), this was causing the browser to spend a lot of time on every single frame, while, since we are aiming to rendering at 60fps, we are left with less than 16ms per frame; moreover, by further analyzing the timeline, we can see 8 "spikes", which bring to the surface another problem: the directives are running too many times, in fact one can see 8 "charts" loading and drawing, while only 4 are being display on screen. Why is this happening?



Figure 5.3: Overview screen, JavaScript profile in the Timeline

This page was built with two versions in mind, one that should be displayed when the cluster-based scoring plugin is available, and one when it's not. In order to achieve this effect without duplicating too much code, I was simply hiding and showing widgets as needed, programmatically. However, for how Angular's directives work, the hiding/showing action takes place after the basic directive functionalities are loaded (needed DOM is linked to the root element, controller's code has run, and so on). Therefore, exploiting the fact that every HTML template I'd need is going to be precached at load time anyway, I decided to split such a page into its two versions; by doing this, I solved the performance issue, while making the code much more readable and maintainable. It's also worth noting that in Angular it can be possible to lazy load the directives, which would have solved the issue as well, but I decided nonetheless to take the "maintainability path".

Much a different process was fixing the performance issues in the review screen. While the previous one had a lot to do with the JavaScript being run, this one finds its

causes in the way CSS properties are transitioned by the browser's rendering engine. First of all, a little bit of context: in the review screen, the transition between filters bar closed and filters bar open states is animated, provoking the subsequent reduction of both the width and the font-size of the tags' table and document container. As it turns out, this can be cumbersome to render at 60fps with the current technologies, for the way the engine takes care of those properties change; for the sake of this report, what matters is that maximum performance can be obtained only when tinkering with transform and opacity, which, as it currently stands, are the only properties that don't trigger neither the layout nor the painting steps of the rendering process.



Figure 5.4: http://csstriggers.com/ - It offers a quick summary of which CSS properties trigger which phase of the rendering process

Therefore, one might think that the fix would simply consist in animating the transform property of the filters, in order to exploit GPU acceleration; however, this was already the case, as the filters' box was translating into view by means of transform: `translateX(0)`. What was missing was a hint to the browser, telling it that the HTML element containing the filters' box was going to change its transform property, thus allowing it to be layer-promoted. Without going too much into the technicalities of this procedure, whenever a moving element needs to be computed by the compositor thread, it also needs to live within its own layer; in the past, in order to place an element inside a new layer, there was the need to use the transform: `translateZ(0)` hack, which was basically forcing layer-promotion. Today however, it suffices to use will-change: transform, which enables every kind of optimization for that kind of transition by the rendering engine, layer-promotion included.

# Chapter 6

## The back-end: architecture, new features and the legacy problem

During my fifth month as intern, I worked on the back-end side of the project, my tasks being refactoring the existing code, developing new features and integrating it on the front-end.

## 6.1 Architecture

As I step in and take charge of the back-end of the web application for CAM, a valuable and rich code base already exists. Let's quickly review the architecture of the system, by listing its components.



Figure 6.1: CAM Architecture

On the front-end, the UI we discussed so far. On the back-end instead, there exist two separate and independent components:

- the *Annotation server*, it's the core engine of the whole CAM system. It's capable of analyzing, mining and augmenting the documents in input, and it outputs in RDF/XML format
- the *Web Server*, which acts as mediator between the client and the Annotation server, has two primary roles: serving static and dynamic assets to the client, and fetching live analysis data from the Annotation server

Since a basic UI for CAM already existed in the past, a Web server is already set up and most of the work needed to parse the RDF results and build an Object

representation of it in memory is already available, and therefore there's no necessity to rebuild everything from scratch. The problem however is that the new client application is expecting certain input data - some of them totally new - with a given structure, in a JSON format, and so on. My job then is to re-factor the existing data into a more suitable shape, develop and build new features and aggregations on top of the available ones, and integrate such a suite of services on the client, by means of a simple RESTful type of service. I thus decided to create a *mediator* module, sitting between client and server, capable of managing the interactions between the two through an API, which exposes the minimal interface needed to make it work; by doing this, I decouple as much as possible both the UI and the already existing Web server, allowing future versions of either one of the other to be easily plugged in or replaced. In addition to this, my module abstracts the interactions and hides as much as possible the underlying processing jobs that take place on the web server, aiming to support Mondeca's decision to outsource CAM's UI maintenance to another company in the future.

## 6.2   Module's structure

The new module exposes the following services:

- `getAnalyzeData`, which returns all data needed by the Analyze page
- `getOverviewData`, `getReviewData` and `getTroubleshootingData`, which take a document and return the relative data
- `indexDocument`, which takes a document and perform the right analysis, depending on the document's type (is it URL, free text, binary data, and so on)
- `getProgress`, which returns the current progress of the analysis process
- `publish`, which takes all user's modifications made to the analysis results, and stores them

This reduces by a big amount the total surface of the API, which used to be bigger and, therefore, much less maintainable and too tightly coupled to the client's implementation. On the other hand, by using less finely grained services there exist the risk of making it too difficult to split pages on the client in smaller pages. This is a trade-off between ergonomics and adaptability, however I prefer this approach since it allows new developers to easily get on board and start writing code from their first day on both the client and server, without having too much to learn about client-server interactions; in addition to this, offering few and "big" services encourage front-end programmers to fetch all the needed data in one single request at load time, which, as described in Chapter 5, can be a big win in terms of both perceived and computational performance. Even though the basic structure of the module might seem completely RESTful, there are a couple of exceptions:

- the `login` service, which makes use of cookies to authenticate the user in the following requests; this isn't completely RESTless, since it doesn't make use of such cookies to store session data
- the `getProgress` service, which, at every call, exposes some session-tied state the server is keeping (the progress of the analysis asked by the user)

## 6.3   Data refactoring

As mentioned in the previous paragraphs, one of the main goals of the mediator module is to re-factor the result into a more suitable format, while enriching it by pre-computing useful aggregation such as totals, counts, and so on. In other words, starting from a list of tags, I moved the complexity on the server by preparing all the statistics needed in the Overview page, collecting the list the Dropbox files in the configured folder for the Analyze data (anticipating the eventual request), and so on. In addition to this, the module enriches the output of the analysis by flagging the annotations basing on their category (is a given tag an inference or has it been extracted from the text?). The whole process might seem a small addition to the global picture, however it is of the uttermost importance that the Web Server undertakes the task of preparing all the data the client might need, supporting as many activities, aggregations and flexibility as possible, and this stands true especially if one thinks of the API as an abstraction layer, enabling client applications that are insightful and performant.

## 6.4   New features and legacy's improvement

Because the new *CamAPI* I introduce sits in between client and server, I had the possibility to rethink the way the two endpoint interact with each other, focusing on improving both performance and user experience; even though the legacy code plays the biggest part of the back-end implementation, there are plenty of opportunities to enhance the existent. Therefore, I contributed with four new features: asynchronous indexation, parallelization, results caching and progress tracking.

**Asynchronous indexation**

In the old CAM, while a document is being processed, the client waits for the process to be finished and gets the result back as soon as it completes in a synchronous fashion. There are two severe drawbacks of this approach: on one side, the client is forced to process one resource at a time, on the other side, the full result of the analysis needs to be transfered on the wire, which may be fairly big and caused in the past several problems for the too heavy payload. My solution for this issue is to break the interaction into smaller pieces and make the exchange of information between client and server asynchronous, as follows:

- client initiates indexation process of a document
- server processes the result
- client is free to ask for the result as soon as it needs it

Such an approach is not only more suitable for an user interface (which usually implies a lot of asynchronous operations), but opens up to new opportunities.

**Parallelization**

The most important enhancement made possible by the asynchronism of the requests is the capability for the client to trigger more than one analysis at once, in a parallel fashion. In my new design, indeed, a "pool" of documents is created and it is then sent to the server for analysis. The advantage is evident: there is now a much higher throughput of processed documents. From an implementation point of view, one could have done better than I did; in fact, in order to process more than one document at the same time, I simply open more than one connection to the server (the number of which is, in most browsers, limited to 6). The web server handles every incoming request with a separate thread, thus true parallelism is achieved. However, a much smarter solution would have been sending the list of documents to the server (the so called *pool*), and have it creating threads; in this way, more than 6 documents could be processed at a time, and just a single network request should have been opened on the client. The reason why I didn't take this path is for it would have taken too much time for me to implement a thread pool on the server, and I was told not to spend too much energy on the back-end, since the company is mostly interested in an usable interface, rather than having me prioritizing parallel execution of more than 6 resources. In addition to this, when the pool contains a lot of documents, the whole system is blocked waiting for every resource - possibly big files - to be transferred to the server, therefore canceling the benefit of parallelism.

**Results caching**

In order to enable the client to asynchronously request for the result of a given document's analysis, it is necessary to have a session-linked results cache on the server; this ultimately solves the problem of having a possibly too big payload to be served in a single connection, since the client can now specifically ask for the parts of the result it is interested in at that precise time: when the user is on the Overview page, there's no need to fetch the data that would only be used in the Review page, nor the log messages that can be accessed in the Troubleshooting section. Since in the desktop environment CAM targets the network is almost always the performance bottleneck, such a strategy has the more general benefit of increasing the overall responsiveness of the application.

**Progress tracking**

Finally, I introduce in CAM a progress bar showing the percentage of completion of every document (which makes sense now that parallel execution is present), plus a global progress indicator which relates to all the resources being processed. As I proceeded submitting my plan to the members of the team who used to work on CAM in the past, I immediately encountered some resistance in introducing the progress tracking; this was not because there is something wrong with the idea itself, but it turned out there's no actual way to get such information from the annotation server, which basically processes the whole document in a stateless manner.

While this may seem a small detail in the entirety of the CAM application, there are some considerations to make on the user experience of it when no progress tracking is available. One of the most important phase of the interaction between a person and an object (being it a real or a virtual one) is the **feedback** the person is presented with, indeed, without some sort of hint that the interaction took place, we usually tend to be doubtful of the success of the operation: *did I click the button? did the system receive my command?* While this is very true for every interaction, it is even more fundamental when talking long activities. In the old CAM there is no way to estimate the amount of time remaining, so the user can't take decisions upon that: *can I switch to some other task while I wait? should I skip that resource, and work on simpler ones for the moment? Am I to wait an hour or a minute?* Thus, I started to investigate how this may be achieved, because I believe that a rough and imprecise indication is much better than no indication in this case.

By digging a little bit more into the problem, I found out that, while the annotation server does its job in a single run, the web server still needs to post-process the result in several ways, before storing the result. In such a multi-step way of proceeding I saw an opportunity to achieve my goal: I identified 5 steps that measured a non-negligible amount of time to complete and used those as an indicator of how far in the processing of the document the system is. I am highly confident that, even though the 5 steps don't take the same time and they are really not much meaningful in terms of the actual activities they correspond to, it is still a big win to be able to present to the user of an estimation of how much work is done and how much is yet to be.

# Chapter 7

## *Wrap up: i18n, future work and conclusions*

In order to conclude the description of the CAM project, it is worth mentioning how the application is internationalized; in addition, since Mondeca believes in the future of such a product, some considerations on the future work that can be done to further improve the experience are included in this final chapter.

## 7.1  i18n

A special effort required on my side has been internationalizing CAM in order to enable translation of the app at two different levels: the goal is to support multiple languages in both the UI's text and labels, and in the underlying data; for instance, a French document can be then reviewed with in English, meaning the labels for classes, entities and annotations will be displayed in the selected language.

### 7.1.1  The UI level

On the UI, my approach consists in relying on a configuration file, more precisely a JSON structure, containing all the labels that are used in the interface; the configuration allows to plug in as many translations as needed, under the form of additional JSON structures. From a technical point of view, I created a *Translations* module to set up all the supported languages; in addition, the *translate* Angular's filter is used in order to enable dynamic change of language, via a simple pop-up window. What this means is that, at any point in time, the user can select another language and see the whole application translated in almost no time. The benefits of the selected solution are the possibility for a non-technical member of the company to review, and easily change, the copywriting, or to an external translator to quickly add a new language by simply providing the key/value pairs (e.g. 'HELLO': 'Salut'), which can be very important when a company needs to move fast and to adapt their product to the customer's needs.

### 7.1.2  The data level

As already mentioned, the underlying system can be tuned to analyze resources in different languages; in addition to this, the taxonomy can, and usually does, contain labels in different languages: the *House* entity therefore will have a `label_en` *House*, a

`label_fr` *Maison* and a `label_de` *Haus*, depending on the configuration. By exploiting this information, the web application allows the user to change the language for the data, by choosing among the available ones advertised by the annotation server. The implementation is pretty straightforward: the web server lists the available languages at bootstrap time, then sends every label in multiple languages in order to avoid having the client requesting the data twice, the UI displays data in the selected (or default) language. Therefore, I simply introduced a *i18n* module, which handles default language, current language, corner cases such as *no language* and provides a simple `extractName` method to be used by the views to extract the label in the correct language from the given entity or class or else.

## 7.2   Future work

## 7.3   Conclusions

# *Bibliography*

Cao, Jerry, Tom Green, and Marek Bowers. 2015. *The Guide to Interactive Wireframing.* UXPin.

Cao, Jerry, Kamil Zieba, Krzysztof Stryjewski, and Matt Ellis. 2015. *Web UI Design for the Human Eye: Colors, Space, Contrast.* UXPin.

Colborne, Giles. 2010. *Simple and Usable: Web, Mobile and Interaction Design.* New Riders.

Nielsen, Jakob. 1993. "Response Times: the 3 Important Limits." In *Usability Engineering.* Morgan Kaufmann Pub.

Norman, Donald. 2013. *The Design of Everyday Things: Revised and Expanded Edition.* Basic Books.

Soegaard, Mads, and Rikke Friis Dam. 2015. *The Encyclopedia of Human-Computer Interaction: 2nd Ed.* https://www.interaction-design.org/literature/book/the-encyclopedia-of-human-computer-interaction-2nd-ed.